



# Introduction to Database Systems

## I2DBS – Spring 2023

- Week 8:
- Indexes

---

Jorge-Arnulfo Quiané-Ruiz

Readings:  
PDBM 12

# Information

- **Homework 2:**

- We will provide feedback and solution ASAP

- **Exercise 8 is online**

- It is about impact of indexes
- = Topic for today (and partly next week)

- **Homework 3**

- It will be online early next week (or before)
- Yet, you have 2 working weeks to solve it: **deadline on April 21<sup>st</sup>**
- Topic = DDL + Normalization + Indexing (+ SQL)
  - Rather extensive ... but you need to learn all of this!
  - Exercise 8 is very useful as preparation!



# *Profile of the Week*

# Edward M. McCreight

## **Co-Inventor of B-Tree Data Structure**

- **1940:** USA
- **1968:** co-invented B-Tree  
"A Space-Economical Trie Storage Structure"
- **1969:** PhD in Computer Science from Carnegie Mellon
- **Institutions:** Boeing, Xerox, Adobe





*Maintain  
& tune*



**RDBMS**

Database Administrator  
(DBA)

# *Indexes*

---

**Readings:**  
PDBM 12

# Initial Case Study

- **Which of these queries should run faster?  
How much faster?**

(1) `SELECT COUNT(*) FROM movie  
WHERE year=1948;`

(2) `SELECT COUNT(*) FROM movie  
WHERE year=1920 or year=1924 or year=1928 or year=1932  
or year=1936 or year=1940 or year=1944 or year=1948;`

**(Q1) If ran several times, would the first run be slower? Why?**

**(Q2) If so, how big would be that difference?**

**(Q3) Would an index improve performance?**

# Seeing Whether an Index is Used

- **EXPLAIN ANALYZE** can be used to show PostgreSQL's query plan

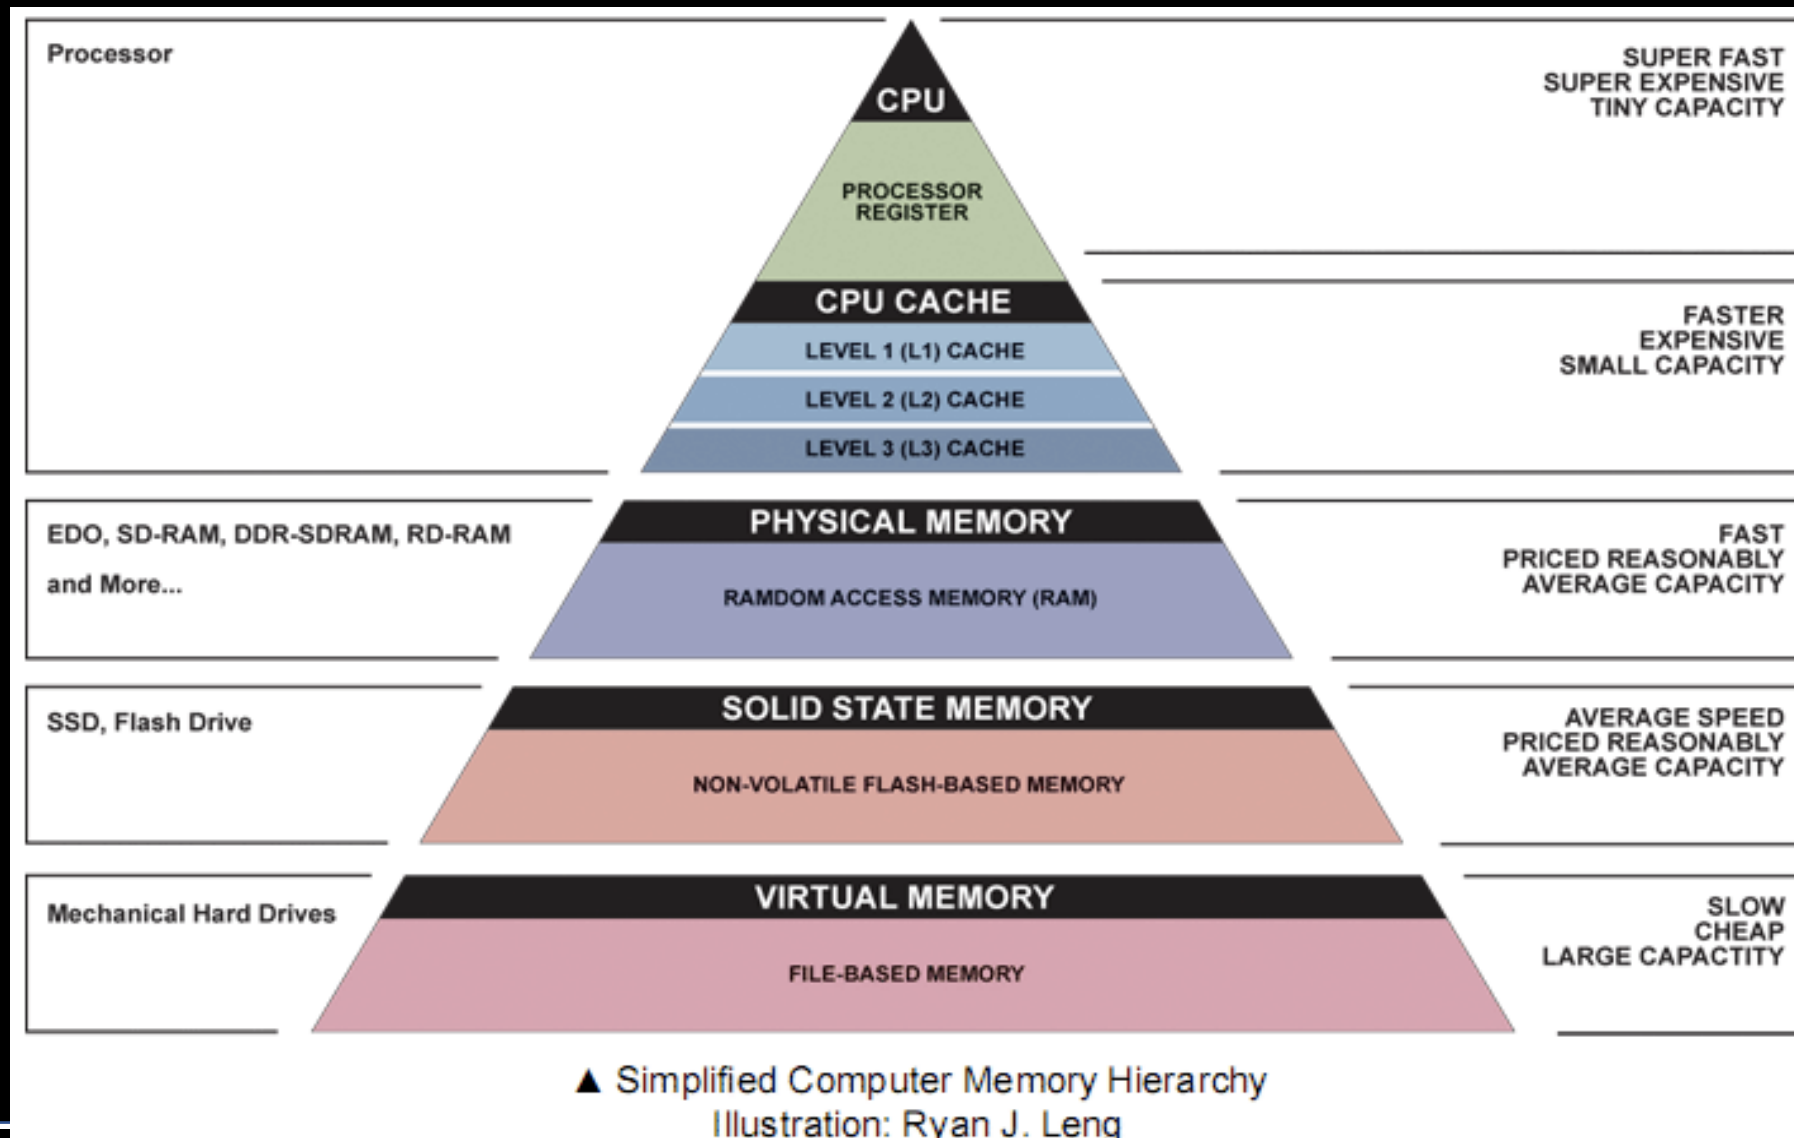
(1) **EXPLAIN ANALYZE**

```
SELECT COUNT(*) FROM movie  
WHERE year=1948;
```

(2) **EXPLAIN ANALYZE**

```
SELECT COUNT(*) FROM movie  
WHERE year=1920 or year=1924 or year=1928 or year=1932  
or year=1936 or year=1940 or year=1944 or year=1948;
```

# Disk Recap





# Impact of Disk-Based Storage

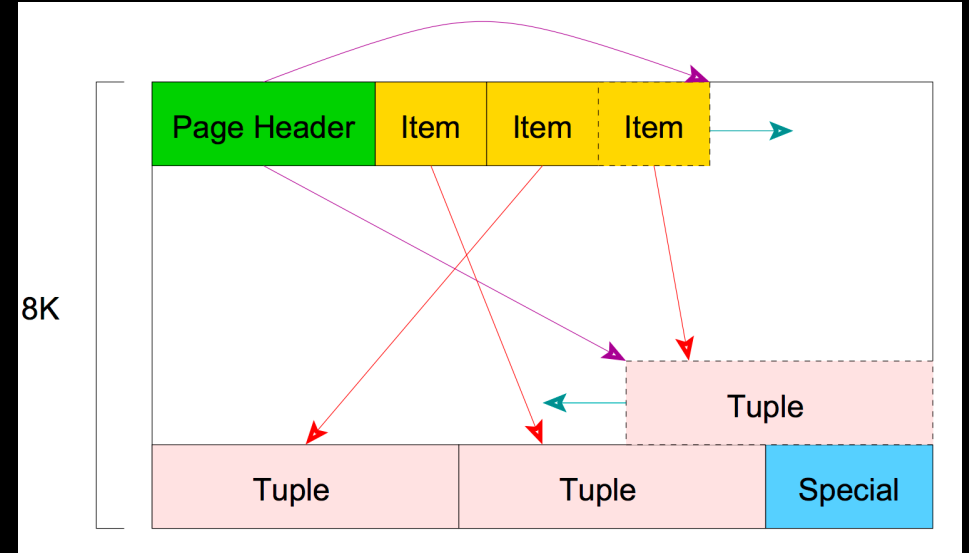
- **Unit of disk reads: KBs**

- We cannot read a few bytes from disk!
- Traditional DBMS: 8KB / 16KB
- Linux: 128KB
- How many employee records fit in one disk read?

- **Historically:**

Avoid reading in random order

- Penalty 1: Random reads costly (but SSDs!)
- Penalty 2: Have to read the same page often!



# Random or Sequential?

## ● Example:

- Table = 2B rows = 1 TB
- Memory = 256 GB  $\Rightarrow$  only small part of the table fits in RAM!
- Disk read = 128 KB  $\Rightarrow$  8M sequential reads vs 2B random reads!
- Cost of reading records in sequential or random order?

## ● Costs:

- HDD:
  - Sequential = ~8M SRs = 1.5 hours
  - Random = ~2B RRs = 114 days
- SSD:
  - Sequential = ~8M SRs = 7 minutes
  - Random = ~2B RRs = 30 hours

	ms/IO	IOs	msec	sec	min	hours	days
HDD	0,642	8.000.000	5136000	5136	86	1,43	0,06
	4,930	2.000.000.000	9860000000	9860000	164.333	2.738,89	114,12
SSD	0,056	8.000.000	448000	448	7	0,12	0,01
	0,055	2.000.000.000	110000000	110000	1.833	30,56	1,27

# Full Table Scans

- When a DBMS sees a query of the form:

```
SELECT *  
FROM R  
WHERE <condition>
```

**It reads through all the tuples of R and report those tuples that satisfy the condition.**

# Selective Queries

- Consider the query from before:

```
SELECT *  
FROM R  
WHERE <condition>
```

- If we have to report **80% of the tuples** in R, it makes sense to do a **full table scan**.
- On the other hand, if the query is **very selective**, and returns just a small percentage of the tuples, **we might hope to do better**.

# Point Queries

- Consider a selection query with a single equality in the condition:

```
SELECT *  
FROM person  
WHERE birthdate='1975-02-06';
```

- This is a **point query**: We look for a single value of birthdate.
  - We may still return  $> 1$  record!
- Point queries are easy if data is sorted by the attribute used in the condition.
  - How? What algorithms would work?

# Range Queries

- Consider a selection query of the form:

```
SELECT *  
FROM person  
WHERE birthdate BETWEEN '1975-02-01' and '1975-02-28';
```

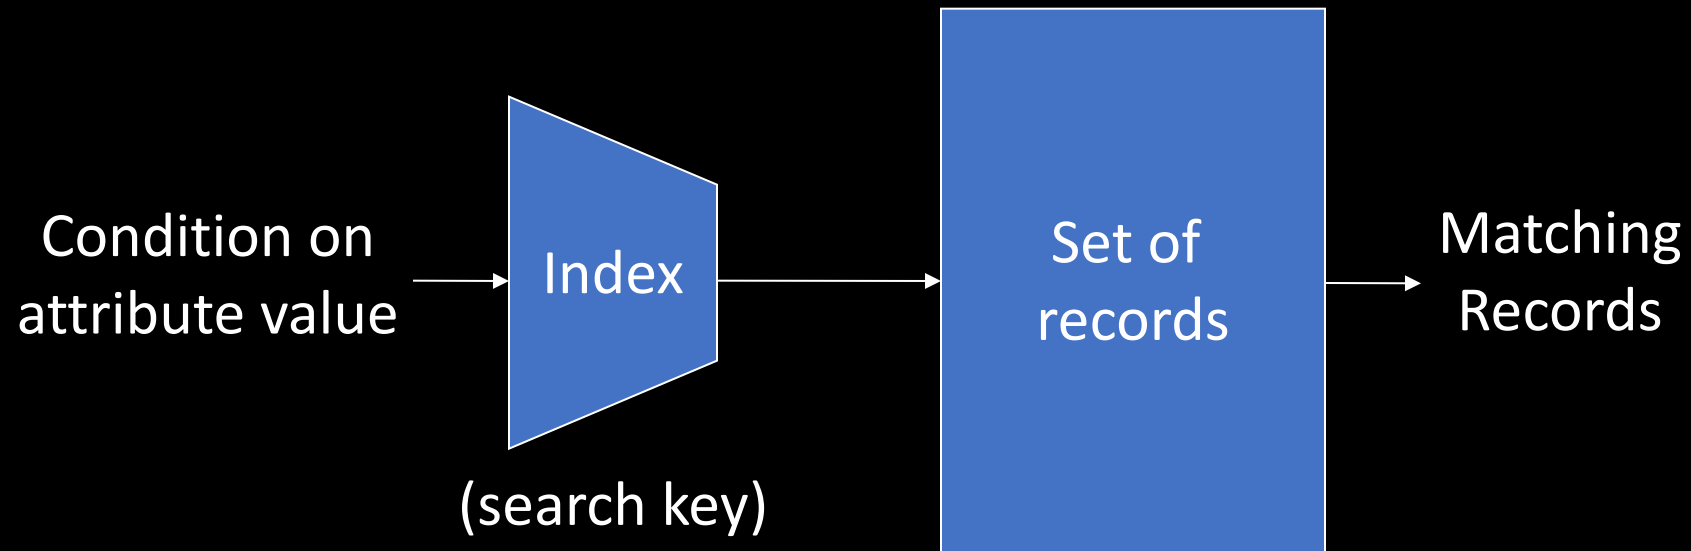
- This is a **range query**: we look for a range of values of birthdate.
- Range queries are also easy if data is sorted by the right attribute.
  - But often not as selective as point queries.

# Indexes

- To speed up queries the DBMS may build an index on the birthdate attribute.
- A database **index** is similar to an index in the back of a book:
  - For every piece of data you might be interested in (e.g., the attribute value 1975-02-06), the index says where to find the row with the actual data!
  - The index itself is organized such that one can quickly do the lookup.
- Looking for information in a relation with the help of an index is called an **index scan** (range) or **index lookup** (point)

# Indexing

- **An index is a data structure that supports efficient access to data**
  - In databases, indexes are also stored on disk





# Two Techniques

---

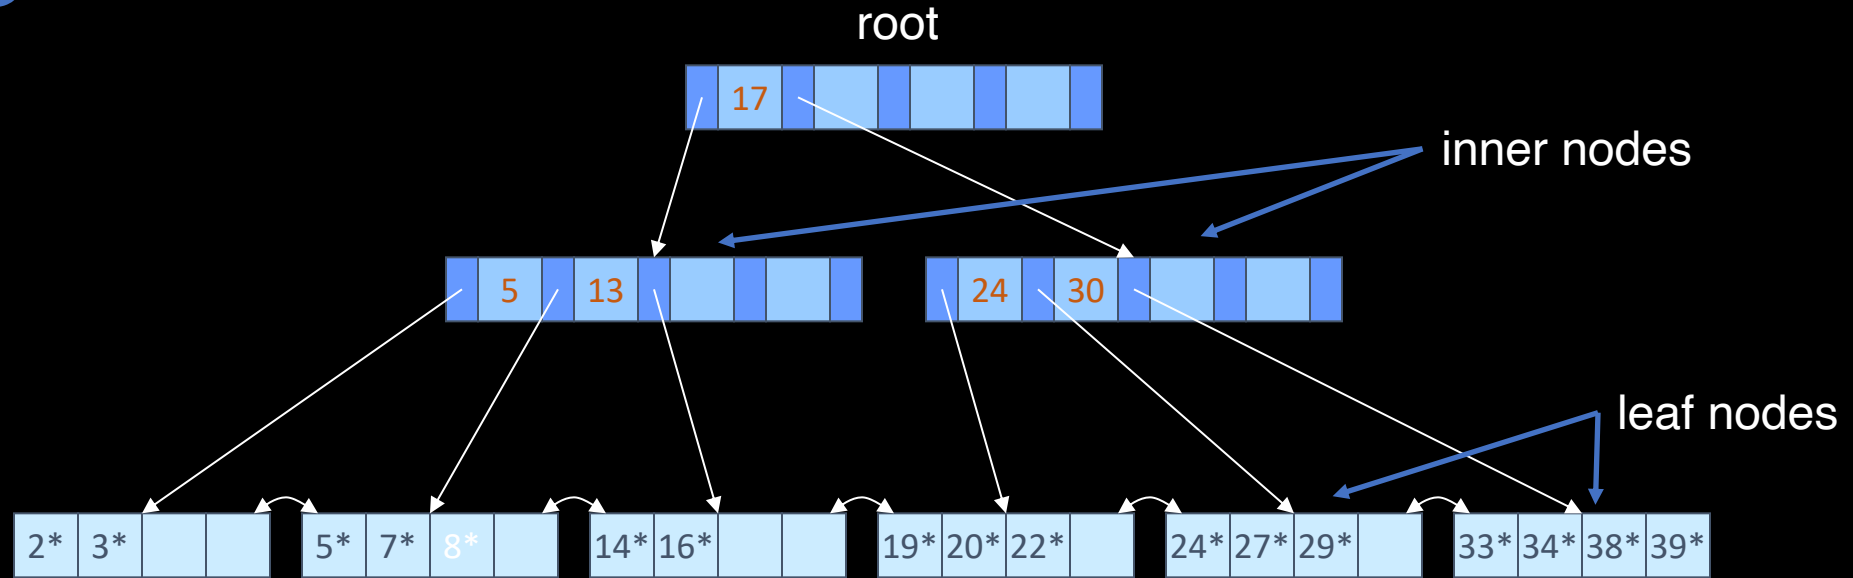
- **Two basic techniques dominate in modern DBMSs:**
  - **Hashing:** Use a fixed transformation algorithm to convert the attribute value into a database address
  - **Tree search:** A dynamic search structure is built that guides the search for a given attribute value to the proper database location
- **Hashing supports equality queries only.**
  - Typically used dynamically for large-scale joins
  - Rarely available to developers
- **Tree search is more versatile and accessible**

# B+ Trees



- **The most common index type**  
... in relational systems
- **Supports equality and range queries**
- **Dynamic structure**
  - Adapts to insertions and deletions
  - Maintains a balanced tree

# A Sample B+-tree



- **X\* represents (search key, pointer list) pairs**  
(X, [address of tuple X1, ...])

- Unique index: Only one entry in list

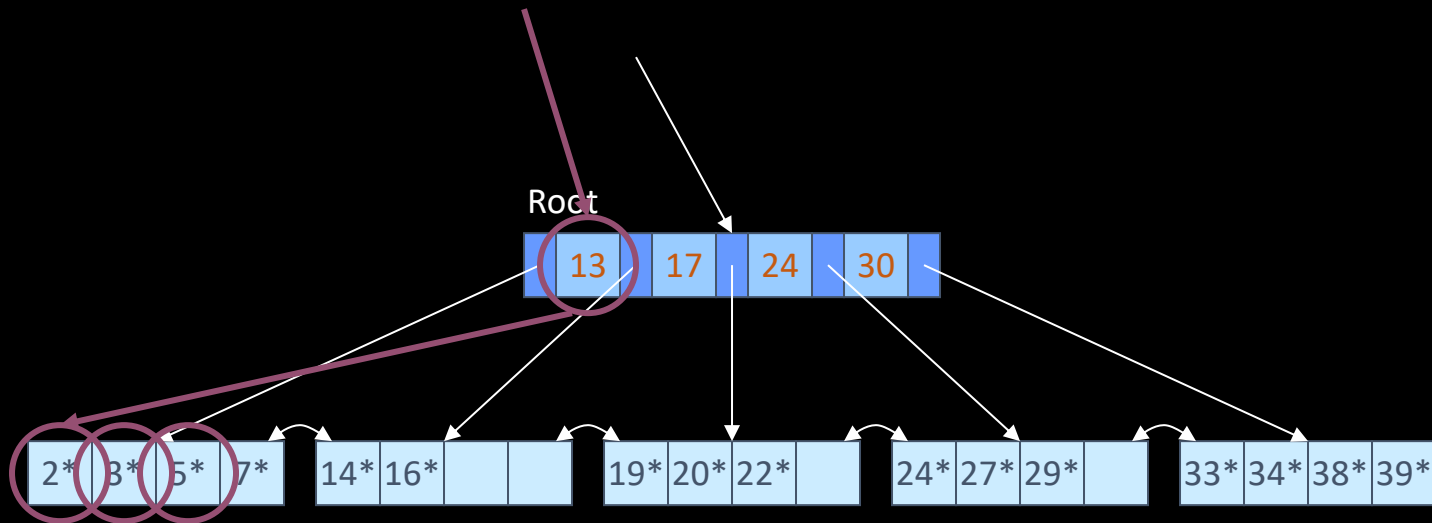
- **Key values are sorted:  $K_1 \leq \dots \leq K_d$**   
(d is maximum capacity or order of node)

- For any two adjacent key values  $K_i, K_{i+1}$  the pointer  $P_i$  points to a node covering all values in the interval  $[K_i, K_{i+1})$

Each inner/leaf node is one disk page  
Nodes have a minimum & maximum capacity

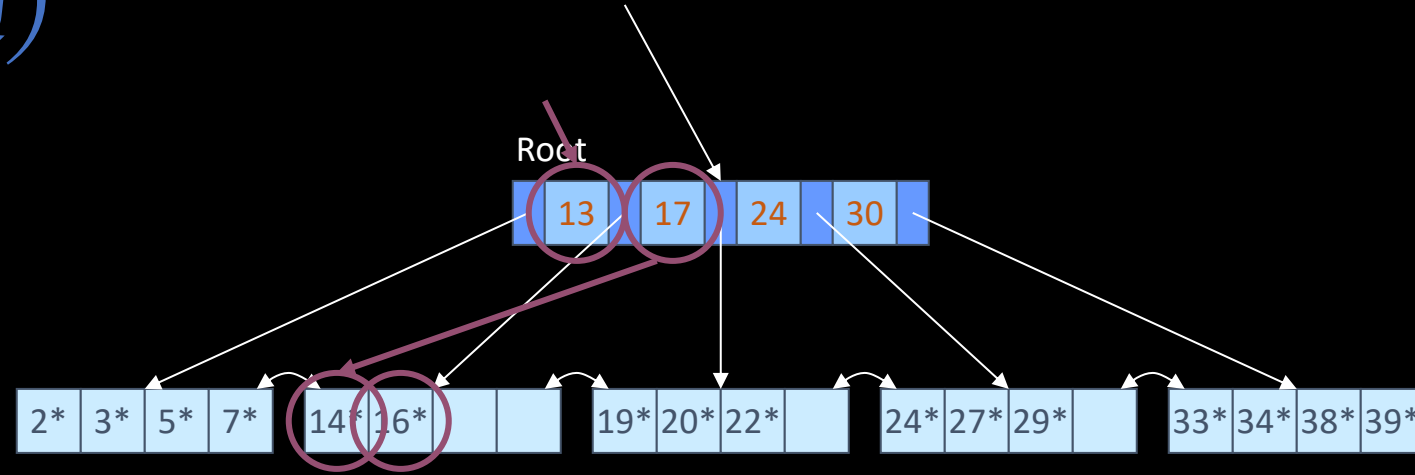
# Searching

- Begin at the root
- Comparisons guide the search to the appropriate leaf
- Ex: Find 5\*

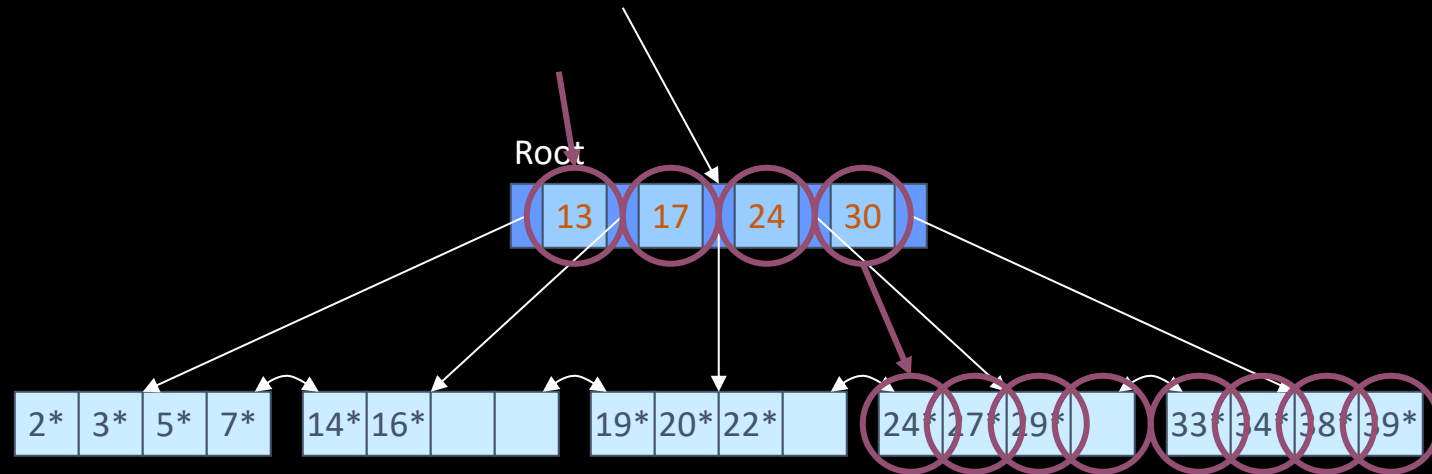


# Searching (11)

- Ex: Find 15\*

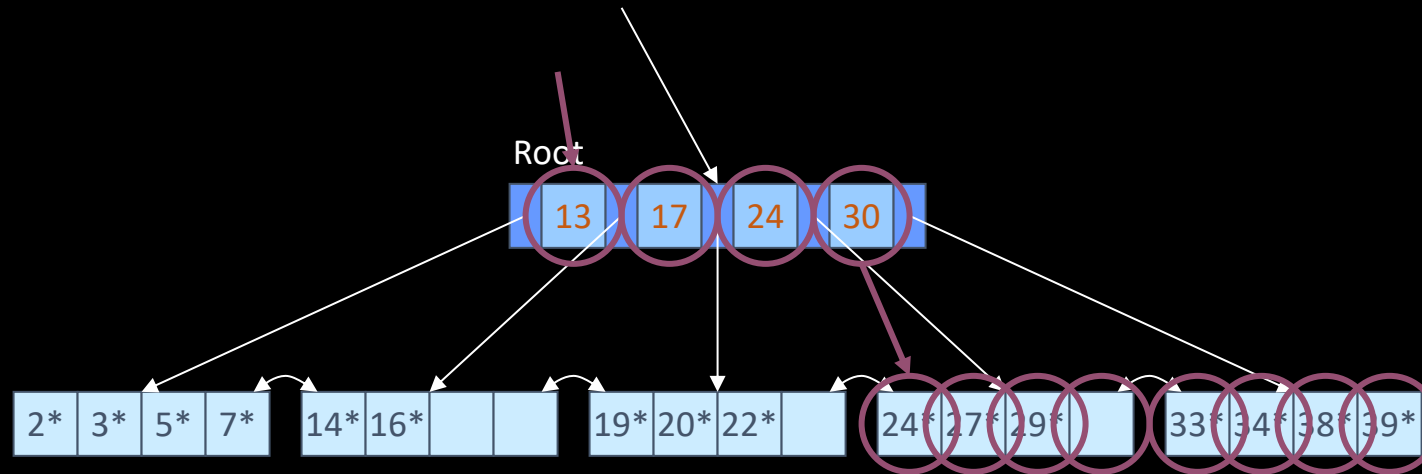


- Ex: Find all records  $\geq 24^*$




# Intra-Node Searching

- We have used scans
- B+-tree nodes have hundred(s) of key values
- Use binary search!




# Inserting

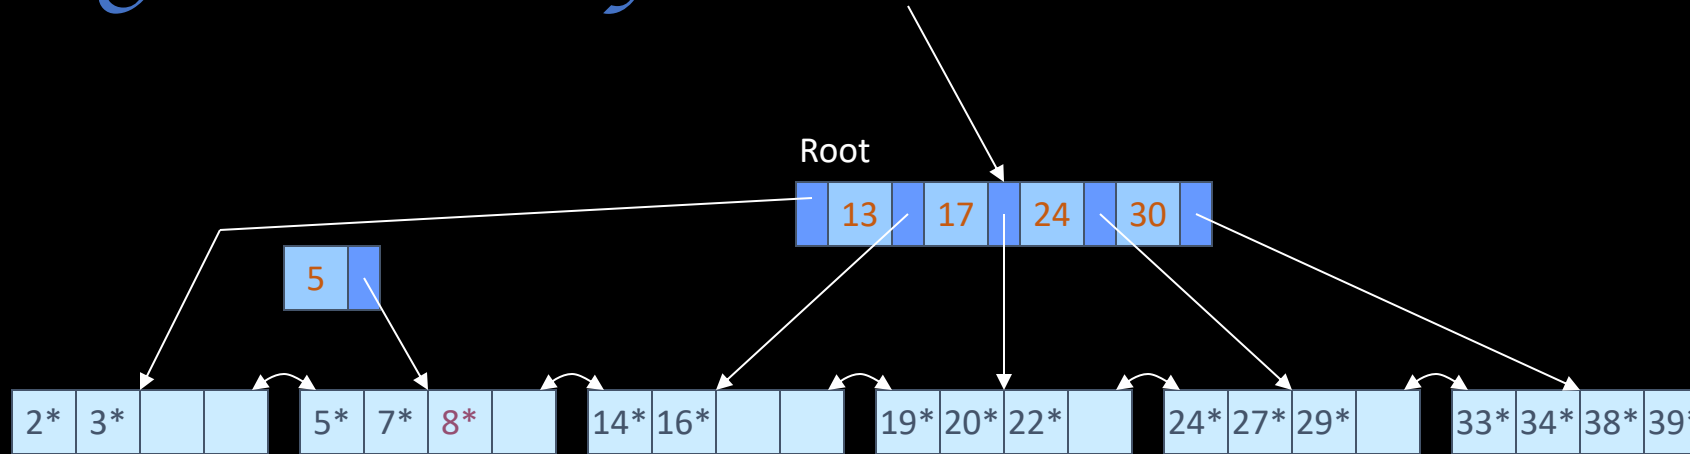
- Leaf is full, must therefore split
- Root is full, must therefore split




- **Inserting 8\* Example**
  - First split the leaf
  - Copy middle search key to the parent



# Inserting 8\* Example (cont.)




- Then split the root
- Move the middle search key into the new root





# After Inserting 8\*

- Trees grow wider, then higher



# Storage Capacity – Some Statistics

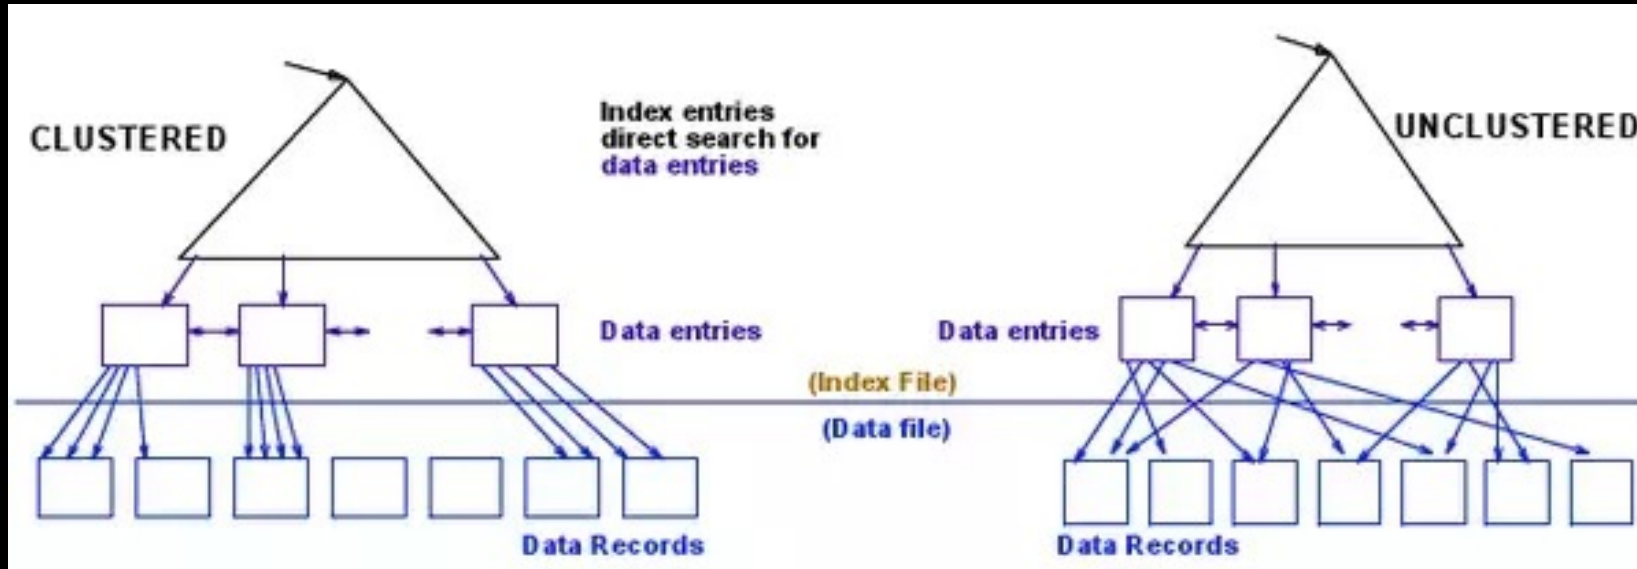
- **A typical tree:**
  - Order: 1000 ( $\sim$  16 KB per page / 16 b per entry)
  - Utilization: 67% (usual numbers in real life)
  - Fanout: 670
- **Capacity**
  - Root: 670 records
  - Two levels:  $670^2 = 448.900$  records
  - Three levels:  $670^3 = 300.763.000$  records
  - Four levels:  $670^4 = 201.511.210.000$  records
- **Top levels may fit in memory**
  - Level 1 = 1 page = 16 KB
  - Level 2 = 670 pages = 11 MB
  - Level 3 = 448.900 pages = 7 GB

# *Index Jargon*

---

- **Indexes vs. Indices**
- **Search key vs. Primary key vs. Candidate key**
- **Unique index vs. Non-unique index**
- **Primary index vs. Secondary index**
- **Dense index vs. Sparse index**
- **Clustered index vs. Unclustered index**

# Clustered vs. Unclustered Index



# Clustered Indexes

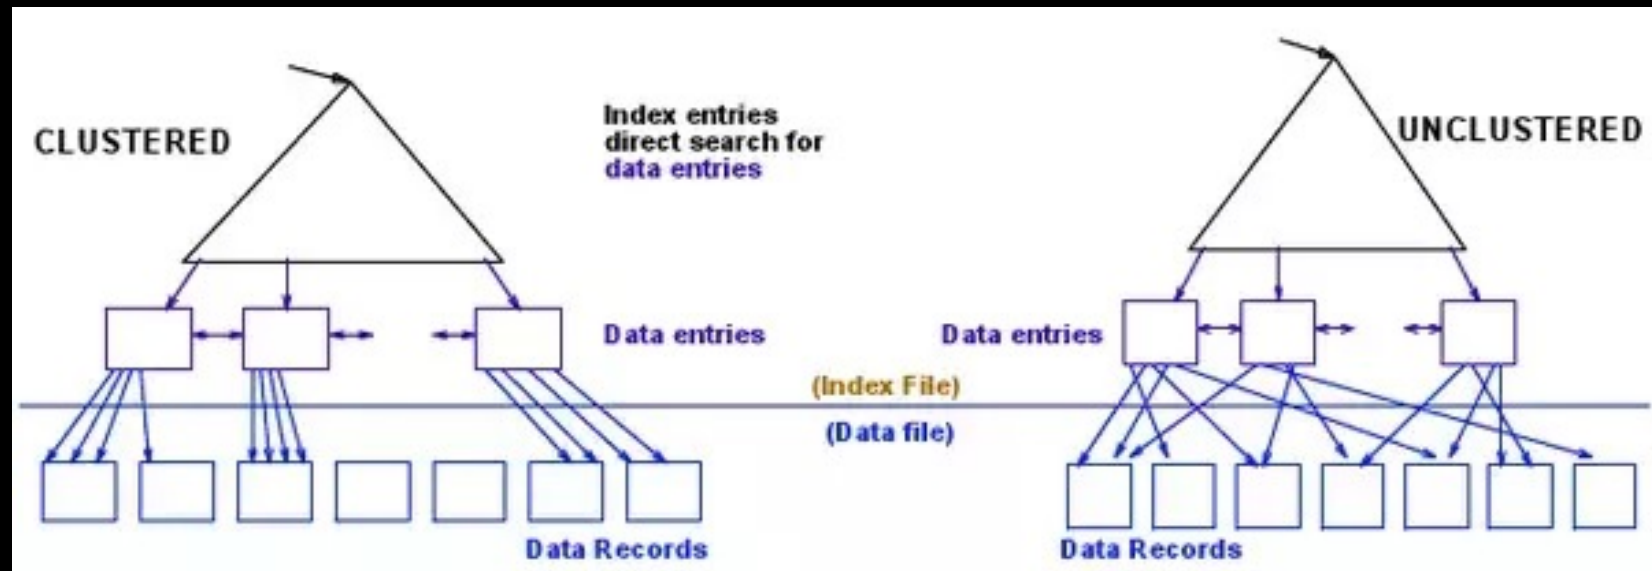
- **If the tuples of a relation are stored sorted according to some attribute, an index on this attribute is called clustered.**
  - Clustered indexes make point and range queries on the key very efficient
  - Why? Sequential reads + As few reads as possible!
- **Many DBMSs automatically build a clustered index on the primary key of each relation.**
  - PostgreSQL has limited clustering support (later!)
- **A clustered index is sometimes referred to as a **primary index**.**
  - Can there be more than one clustered? Why?

# Unclustered Indexes

- It is possible to create further indexes on a relation. Typical syntax:
  - `CREATE INDEX myIndex ON involved(actorId);`
- The unclustered indexes are sometimes called non-clustered or **secondary indexes**.
- **Unclustered indexes:**
  - Make most point queries more efficient.
  - Make some (narrow) range queries more efficient.

# Clustered vs. Unclustered Index

- **To retrieve M records, where M is small:**
  - Clustered: Probably one disk read
  - Unclustered: Probably M random disk reads
- **To retrieve M records, where M is large:**
  - Clustered: Probably  $M/\text{records\_per\_page}$  sequential disk reads
  - Unclustered: Up to M random disk reads
- **We still need to read the index itself – same for both!**




# *Index Scan vs Full Table Scan*

---

- **Point and range queries on the attribute(s) of the clustered index are almost always best performed using an index scan.**
- **Non-clustered indexes should only be used with high selectivity queries.**
  - **Old rule of thumb:** a secondary index scan is faster than a full table scan for queries returning less than 1% of a relation.
  - New rule of thumb?



# Impact of Clustering on Performance



# Covering Index

- An index that contains **ALL** attributes used in a query is called **covering**
  - Resulting query plans are index-only

```
SELECT COUNT(*) FROM movie WHERE year=1948;  
CREATE INDEX movieyear ON movie(year);
```

```
SELECT name FROM person WHERE height=170;  
CREATE INDEX phn ON person(height, name);
```

- The data from the relation is not needed = no disk reads required to retrieve tuples
  - Should a **covering** index be **clustered** or **unclustered**?
  - What is what matters then?

# Clustered vs. Covering Index

- **To retrieve M records, where M is small:**
  - Clustered: Probably one disk read
  - Covering: Definitely need 0 disk reads
- **To retrieve M records, where M is large:**
  - Clustered: Probably  $M/\text{records\_per\_page}$  sequential disk reads
  - Covering: Definitely need 0 disk reads
- **We still need to read the index itself – same for both!**

# Practice: Types of Indexes

- **For each of the following queries:**
  - Which index would give the best plan?
  - Would the index be covering?
  - Would you prefer clustered or unclustered index?
  - Based on these queries, if you could choose, which index should be clustered?

(Q1) SELECT \* ... WHERE birthdate = '20-02-2002'

(Q2) SELECT \* ... WHERE height < 170

(Q3) SELECT \* ... WHERE ID = 4564

(Q4) SELECT AVG(birthdate) ...

# Multi-Attribute Indexes

- **Defining an index on several attributes:**


```
CREATE INDEX myIndex  
ON person (height,birthdate);
```

- **Speeds up point/range queries such as:**

```
SELECT *  
FROM person  
WHERE height=170 and birthdate<'1945-08-08'
```

- **An index on several attributes usually gives index for any prefix of these attributes, due to lexicographic sorting.**

# Lexicographic Sorting?



age, salary

age = 21 22 23 ... 67

salary = 10/11 / ... 90 9/13 / ... 100 11 / ... 99 8 / ... 101

# Problem Session

- Which point and range queries are "easy" (equality on a prefix, range on one) when the relation is indexed with this two-attribute index?

```
CREATE INDEX myIndex  
ON person (height, birthdate);
```

(Q1) A range query on height?

(Q2) A range query on birthdate?

(Q3) A point query on birthdate?

(Q4) A point query on birthdate combined with a range query on height?

(Q5) A point query on height combined with a range query on birthdate?

# Choosing to Use an Index

- **The choice of whether to use an index is made by the DBMS for every instance of a query**
  - May depend on query parameters
  - You do not need to take indexes into account when writing queries
- **Estimating selectivity is done using statistics**
  - In PostgreSQL, statistics are gathered by executing statements such as **ANALYZE** involved



# Choosing Columns

- **Candidates for index search keys**
  - Columns in WHERE clauses
  - Columns in GROUP BY clauses
  - Columns in ORDER BY clause
- **Columns that are rarely candidates**
  - Large columns (too much space)
  - Frequently updated columns (too much maintenance)
  - Columns in SELECT clauses (not used to find tuples)
    - ... but see covering indices!

# *What Speaks Against Indexing?*

- **Space usage:**
  - Similar to size of indexed columns (plus pointer)
  - Most space for leaves, less for tree nodes
  - **Not really important!**
- **Time usage for keeping indexes updated under data insertion/change:**
  - Depends on the index architecture
  - **This is important!**

# *Other Impact of Indexes*

**The DBMS may use indexes in other situations than a simple point or range query.**

- **Some joins can be executed using a modest number of index lookups**
  - May be faster than looking at all data
  - But hash-based joins are usually fastest (next lecture!)
- **Some queries may be executed by only looking at the information in the index**
  - Index only query execution plan ("covering index")
  - May need to read much less data.
- **Consistency (checking keys and foreign keys)**

# Practice: Problem Session

- What would be good indexes for this query?

```
SELECT firstNames  
FROM person  
WHERE gender='m'  
AND firstnames LIKE 'Maria%';
```

# *Index types*

---

- **Common:**

- B-trees (point queries, range queries)
- Hash tables (only point queries on the whole search key, but somewhat faster)
- Bitmap indexes (good for "dense" sets)

- **More exotic:**

- Full text indexes (substring searches)
- Spatial indexes (proximity search, 2D range search, multimedia, ...)
- ... and way more!

# *B+-tree Implementations in Some Major DBMS*

## ● **Postgres**

- Cannot specify a clustered index!
- Manual CLUSTER command!
- Is an index on SERIAL clustered?

## ● **SQL Server**

- Table stored in clustered index
- Primary keys can be unclustered
- Indexes maintained dynamically

## ● **DB2**

- Table stored in clustered index
- Explicit command for index reorganization

## ● **Oracle**

- No clustered index until 10g
- Index organized table (unique/clustered)
- Indexes maintained dynamically

## ● **MySQL**

- Primary key is clustered
- Table stored in clustered index
- Indexes maintained dynamically

# Takeaways

- **Large databases need to be equipped with suitable indexes**
  - Need understanding of what indexes might help a given set of queries
  - 'Key' distinction: Clustered vs Unclustered
  - A detailed understanding of various index types is beyond the scope of this course



# What is next?

- 
- **Next Lecture:**
    - Understanding Query Processing